

Ellen White

Email: egwhite518@gmail.com

Cell: 224-619-0445

In [1]:

```
import pandas as pd
import numpy as np
# import dataset: https://www.kaggle.com/datasets/rishikeshkonapure/hr-analytics
dataset = pd.read_csv("HR Analysis/HR-Employee-Attrition.csv")
dataset.head()
```

Out[1]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationalField
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Life Sciences
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences
4	27	No	Travel_Rarely	591	Research & Development	2	1	Life Sciences

5 rows × 35 columns

Pework - Data Cleansing

In [2]:

```
# Identify Unique/Null Values
nulls = pd.DataFrame({
    'Unique':dataset.nunique(),
    'Null':dataset.isna().sum(),
})
print(nulls)
print(len(dataset.columns))
```

	Unique	Null
Age	43	0
Attrition	2	0
BusinessTravel	3	0
DailyRate	886	0
Department	3	0
DistanceFromHome	29	0
Education	5	0
EducationField	6	0
EmployeeCount	1	0
EmployeeNumber	1470	0
EnvironmentSatisfaction	4	0
Gender	2	0
HourlyRate	71	0
JobInvolvement	4	0
JobLevel	5	0

JobRole	9	0
JobSatisfaction	4	0
MaritalStatus	3	0
MonthlyIncome	1349	0
MonthlyRate	1427	0
NumCompaniesWorked	10	0
Over18	1	0
Overtime	2	0
PercentSalaryHike	15	0
PerformanceRating	2	0
RelationshipSatisfaction	4	0
StandardHours	1	0
StockOptionLevel	4	0
TotalWorkingYears	40	0
TrainingTimesLastYear	7	0
WorkLifeBalance	4	0
YearsAtCompany	37	0
YearsInCurrentRole	19	0
YearsSinceLastPromotion	16	0
YearsWithCurrManager	18	0
35		

```
In [3]: # Drop redundant columns that don't differentiate employee data
# No need to check for duplicate employees because the number of unique employees
for col in dataset.columns:
    if len(dataset[col].unique()) == 1:
        print(col)
        dataset.drop(col, inplace=True, axis=1)
print(len(dataset.columns))
```

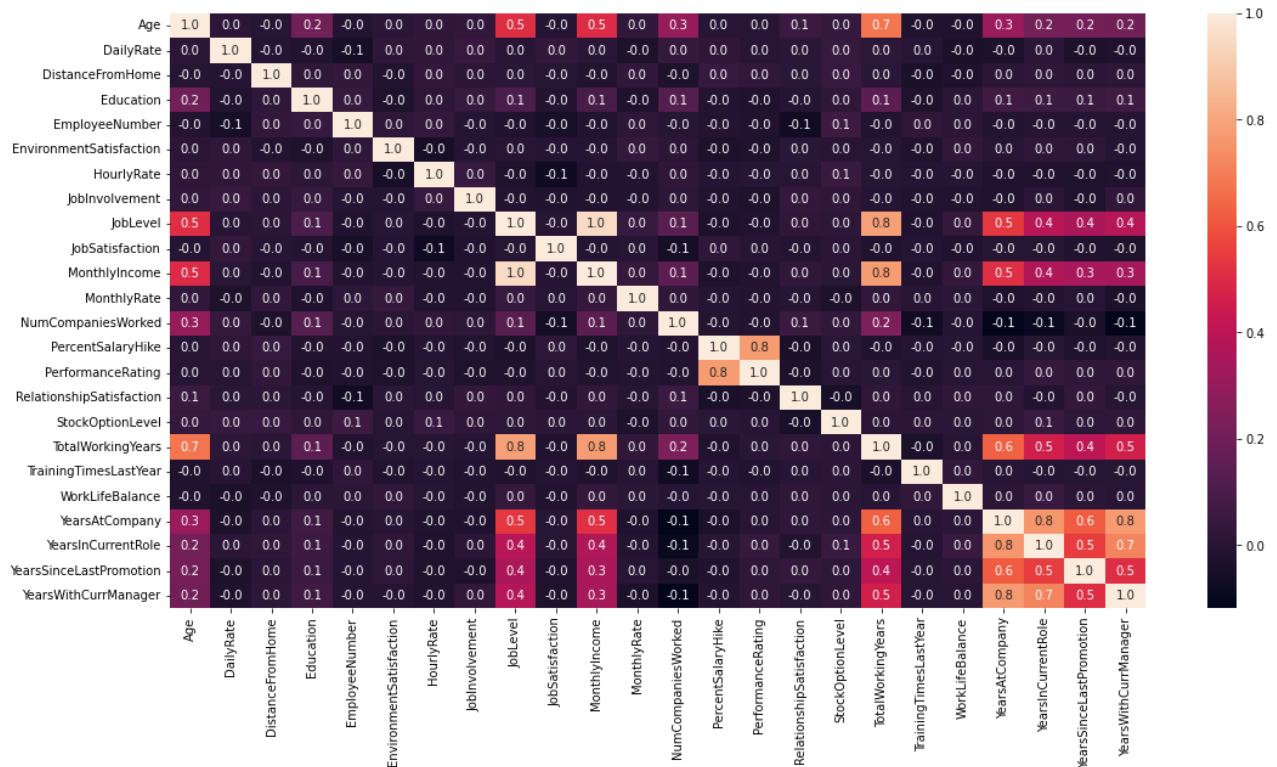
```
EmployeeCount
Over18
StandardHours
32
```

Step 1: Correlation Matrix to determine which variables tend to coincide.

```
In [4]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(18,9))
sns.heatmap(dataset.corr(), annot=True, fmt='.1f')
```

```
Out[4]: <AxesSubplot:>
```



In [5]:

```
# Convert categorical data to numerical data
objects = list(dataset.select_dtypes(include='O'))
print(objects)
numerical = pd.get_dummies(data=dataset, columns=objects)
numerical
```

```
['Attrition', 'BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'OverTime']
```

Out [5]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeNumber	EnvironmentSatisfaction
0	41	1102		1	2	2
1	49	279		8	1	3
2	37	1373		2	2	4
3	33	1392		3	4	4
4	27	591		2	1	1
...
1465	36	884		23	2	3
1466	39	613		6	1	4
1467	27	155		4	3	2
1468	49	1023		2	3	4
1469	34	628		8	3	2

1470 rows x 54 columns

In [6]:

```
plt.figure(figsize=(18,9))
x = numerical.drop(columns = ['EmployeeNumber'])
corrs = x.corr()
np.fill_diagonal(corrs.values, -2)
#corrs
#sns.heatmap(corrs, mask = (np.abs(corrs) >= 0.8))
```

<Figure size 1296x648 with 0 Axes>

In [7]:

```
def get_redundant_pairs(x):
    pairs_to_drop = set()
    cols = x.columns
    for i in range(0, x.shape[1]):
        for j in range(0, i+1):
            pairs_to_drop.add((cols[i], cols[j]))
    return pairs_to_drop

def get_top_correlations(x, n=10):
    au_corr = x.corr().unstack()
    labels_to_drop = get_redundant_pairs(x)
    au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=False)
    return au_corr[0:n]

print("Top Correlations")
print(get_top_correlations(x))
```

Top Correlations		
JobLevel	MonthlyIncome	0.950300
Department_Human Resources	JobRole_Human Resources	0.904983
Department_Sales	JobRole_Sales Executive	0.808869
JobLevel	TotalWorkingYears	0.782208
PercentSalaryHike	PerformanceRating	0.773550
MonthlyIncome	TotalWorkingYears	0.772893
YearsAtCompany	YearsWithCurrManager	0.769212
	YearsInCurrentRole	0.758754
YearsInCurrentRole	YearsWithCurrManager	0.714365
Age	TotalWorkingYears	0.680381
dtype: float64		

The most highly correlated factors are unsurprising, the highest being Job Level and Income. Years and experience contribute to higher pay. You can also see that this company has a common pay for performance philosophy as salary hikes are more correlated with performance as opposed to job level or experience.

Step 2: Gender

In [8]:

```
avgGender = dataset.groupby(['Gender']).mean()
cols = ['MonthlyIncome', 'HourlyRate', 'PercentSalaryHike', 'StockOptionLevel']

fig, axes = plt.subplots(1, len(cols), figsize=(10, 5))

axes[0].bar(avgGender.index, avgGender['MonthlyIncome'])
axes[0].set_title('Avg Monthly Income')

axes[1].bar(avgGender.index, avgGender['HourlyRate'], color = 'red')
axes[1].set_title('Avg Hourly Rate')
```

```

axes[2].bar(avgGender.index, avgGender['PercentSalaryHike'], color = 'green')
axes[2].set_title('Avg Percent Salary Increase')

axes[3].bar(avgGender.index, avgGender['StockOptionLevel'], color = 'orange')
axes[3].set_title('Avg Stock Option Level')

plt.tight_layout()
plt.show()

```



Can conclude that this company doesn't have a pay equity problem to address

Exploring other data

```

In [9]: cols = ['WorkLifeBalance', 'TrainingTimesLastYear', 'PerformanceRating', 'NumCompan',
               'JobInvolvement', 'JobLevel', 'JobSatisfaction',
               'EnvironmentSatisfaction', 'Education']
plt.figure(figsize=(9, 36))
for i, col in enumerate(cols):
    gender_means = dataset.groupby('Gender').mean()

    axes = plt.subplot(len(cols), 2, i + 1)
    sns.histplot(dataset, x=dataset[col], hue="Gender", stat="probability", mult

    ax = axes.twinx()
    male_line = ax.axvline(gender_means.loc['Male', col], color='blue', linestyle=
    female_line = ax.axvline(gender_means.loc['Female', col], color='red', lines

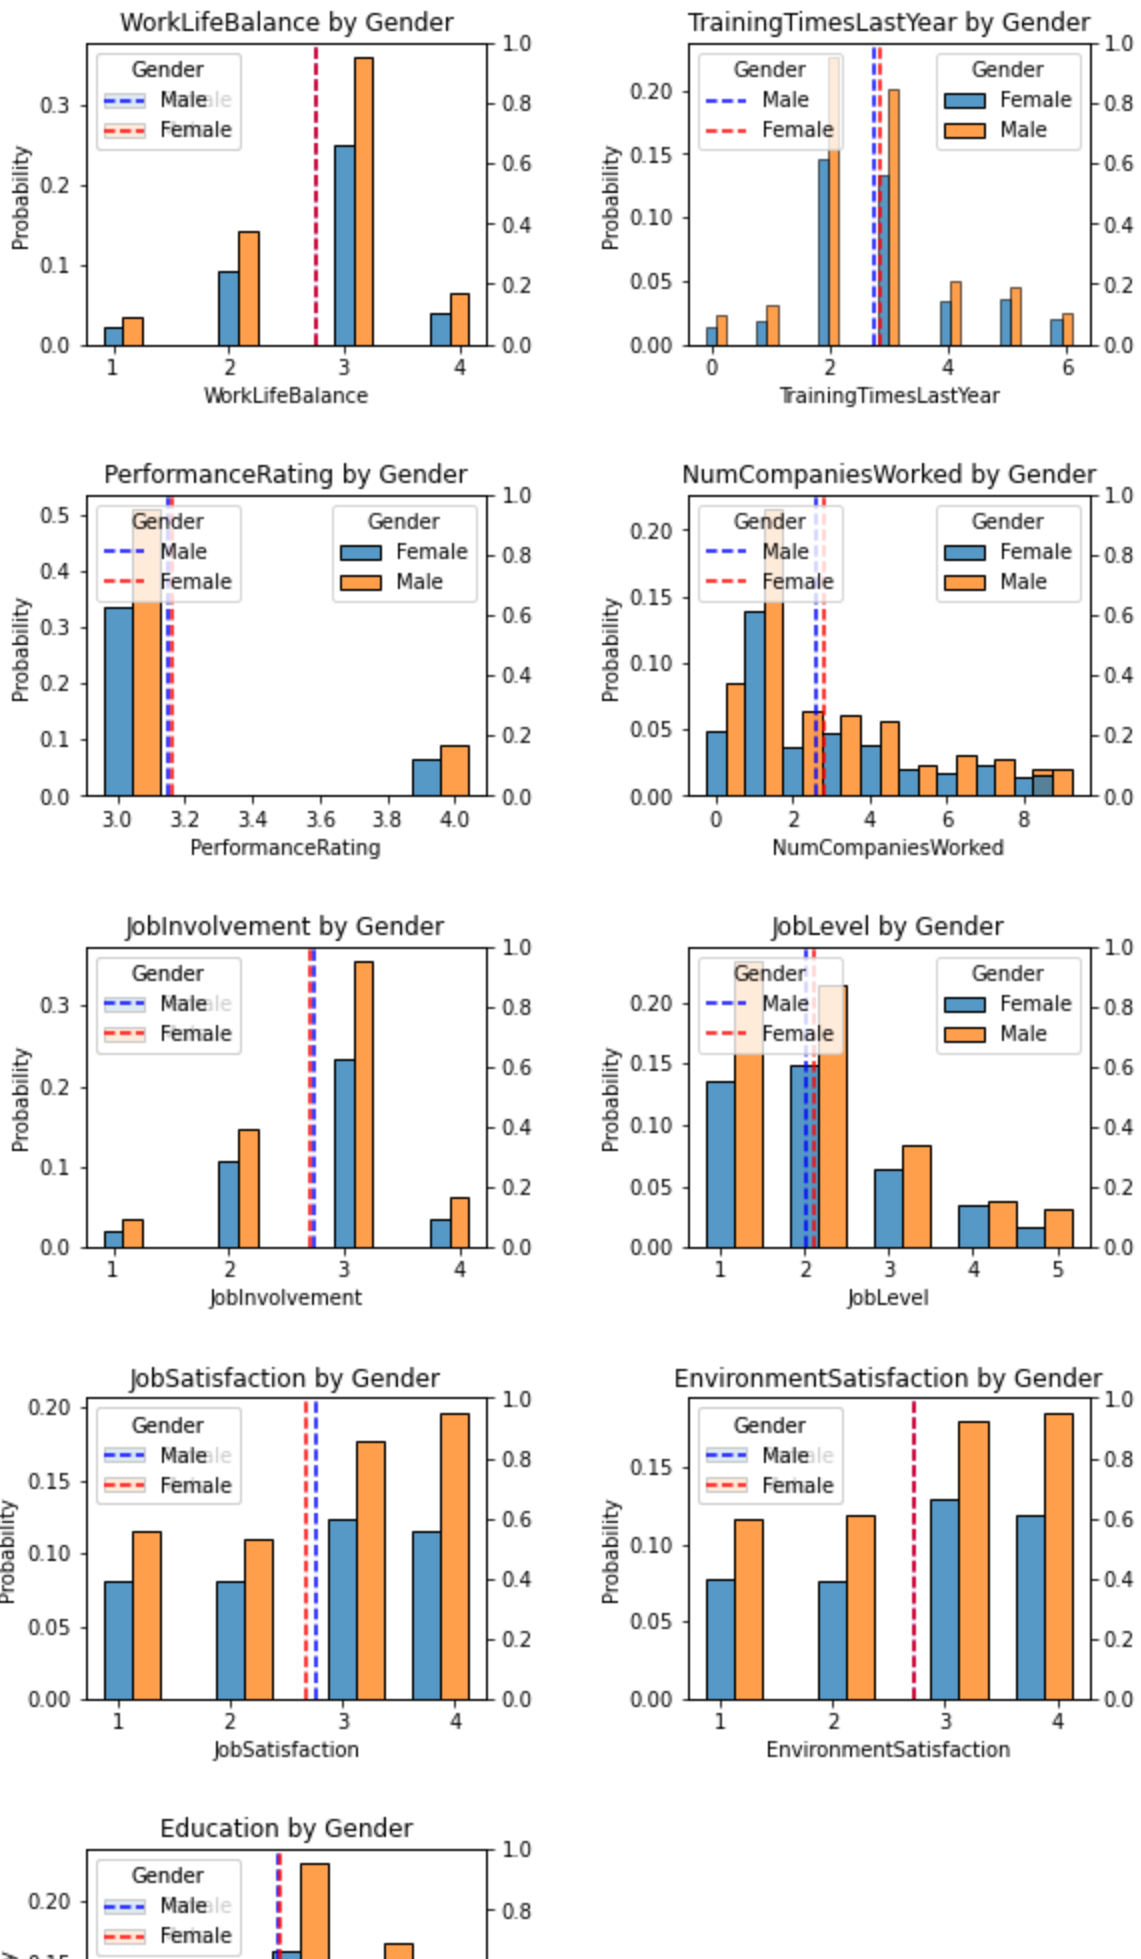
    axes.set_xlabel(col)
    axes.set_title(f'{col} by Gender')


    ax.legend(loc='upper right')
    plt.legend(title='Gender', loc = 'upper left')

plt.subplots_adjust(hspace=.5, wspace = .5)

plt.show()
# ideally figure out how to keep legends from overlapping

```



Most factors are equal, on average, amongst males and females. Males **slightly** tend to be more satisfied with their jobs, while women **slightly** tend to have worked at slightly more companies.



Step 3: Attrition Patterns


Reminder, when attrition = yes, that means the employee left the company.

In [10]:

```
cols = ['Department', 'EducationField', 'Gender', 'BusinessTravel', 'JobRole']

for i in cols:
    fig, axes = plt.subplots(figsize=(16,5))
    sns.barplot(x=dataset[i], y=dataset['MonthlyIncome'], hue=dataset['Attrition'])
    plt.xticks(rotation=50, fontsize=10)
    for cont in axes.containers:
        axes.bar_label(cont, rotation=30, fontsize=10)
    plt.tight_layout()
    plt.show()
```





You can see that income isn't always the driving force of attrition. For example, the research directors who left tended to have higher salaries than the research directors who stayed. This is a case where the company may want to explore issues in the research department that may be driving employees out.

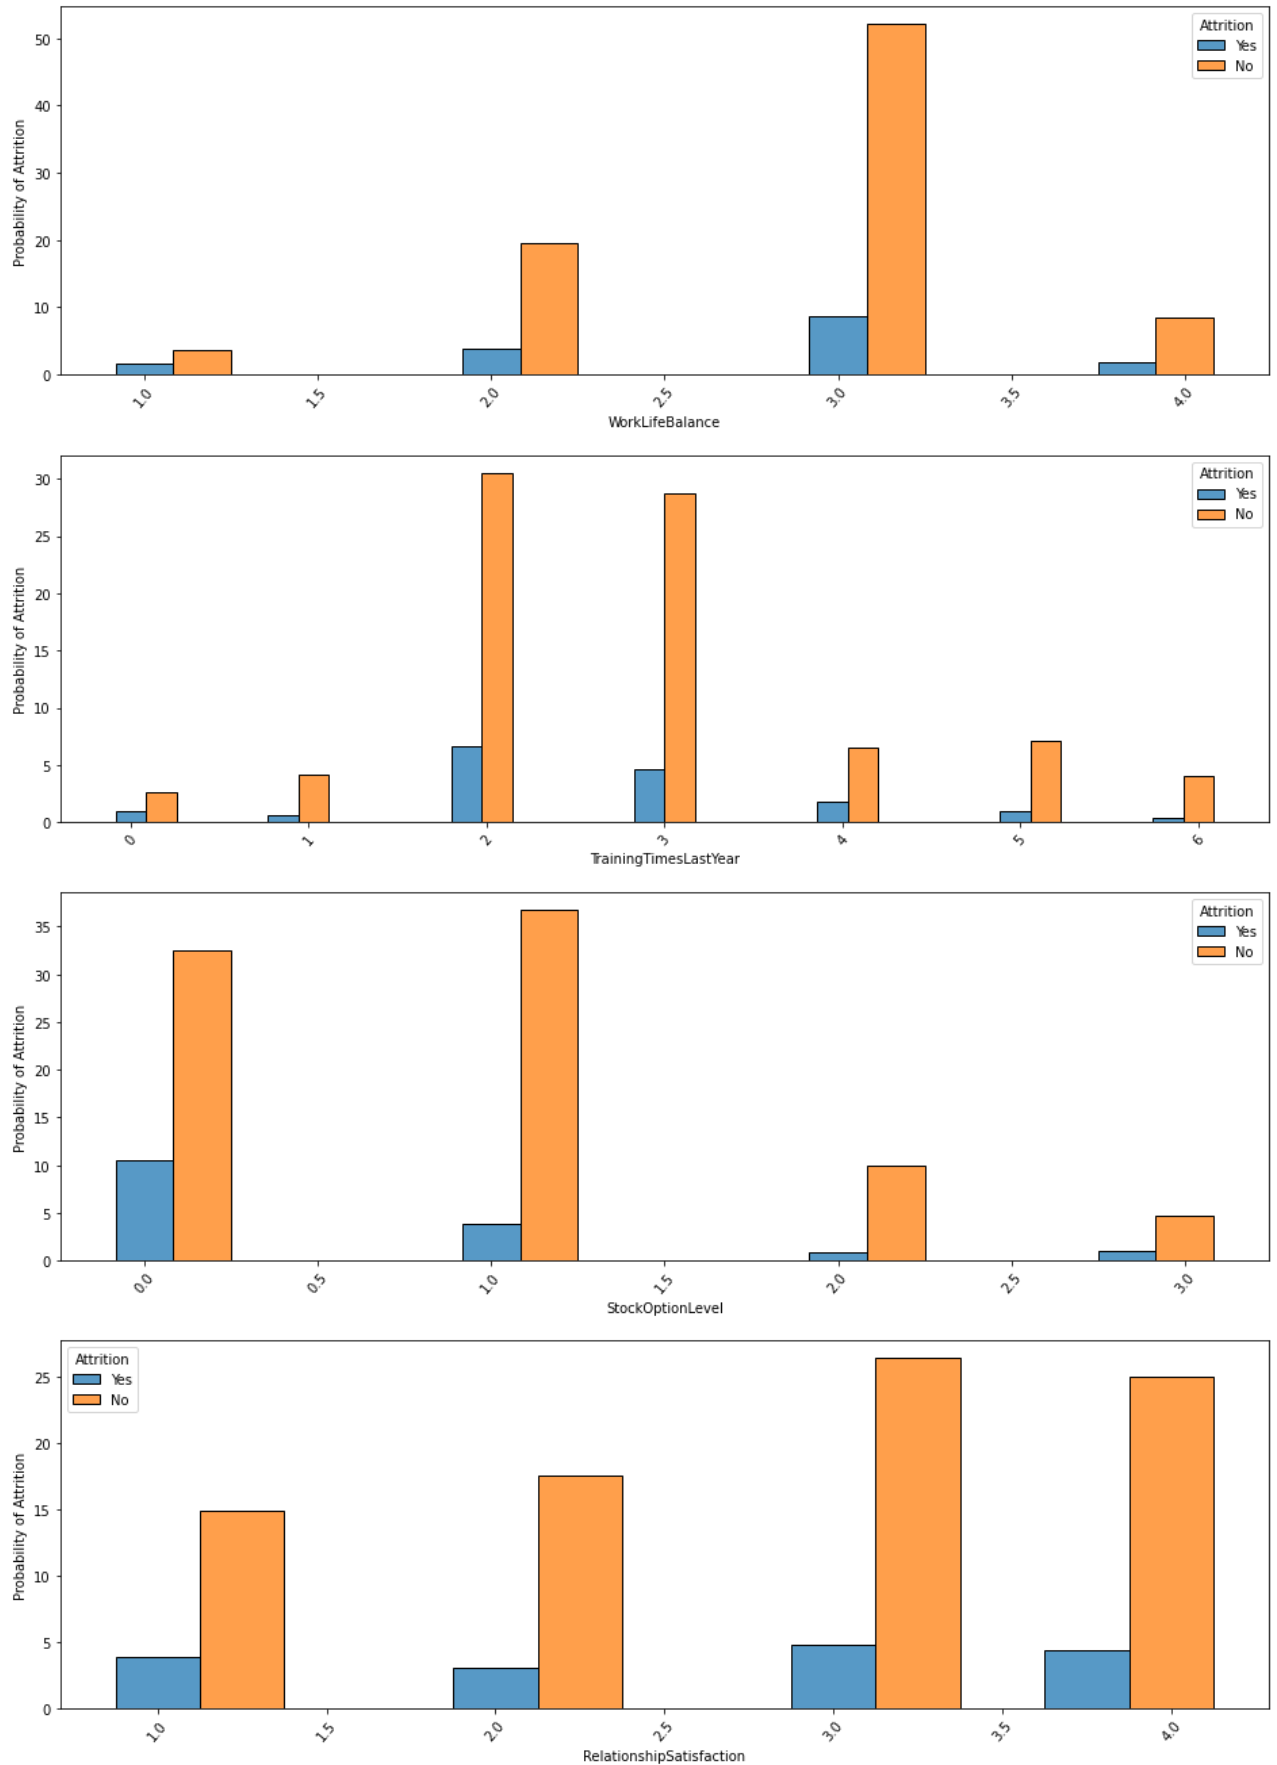
These are some additional factors independent of income. You'll see there's no overwhelming evidence of turnover in any of these cases, even in the case of low environmental/job satisfaction.

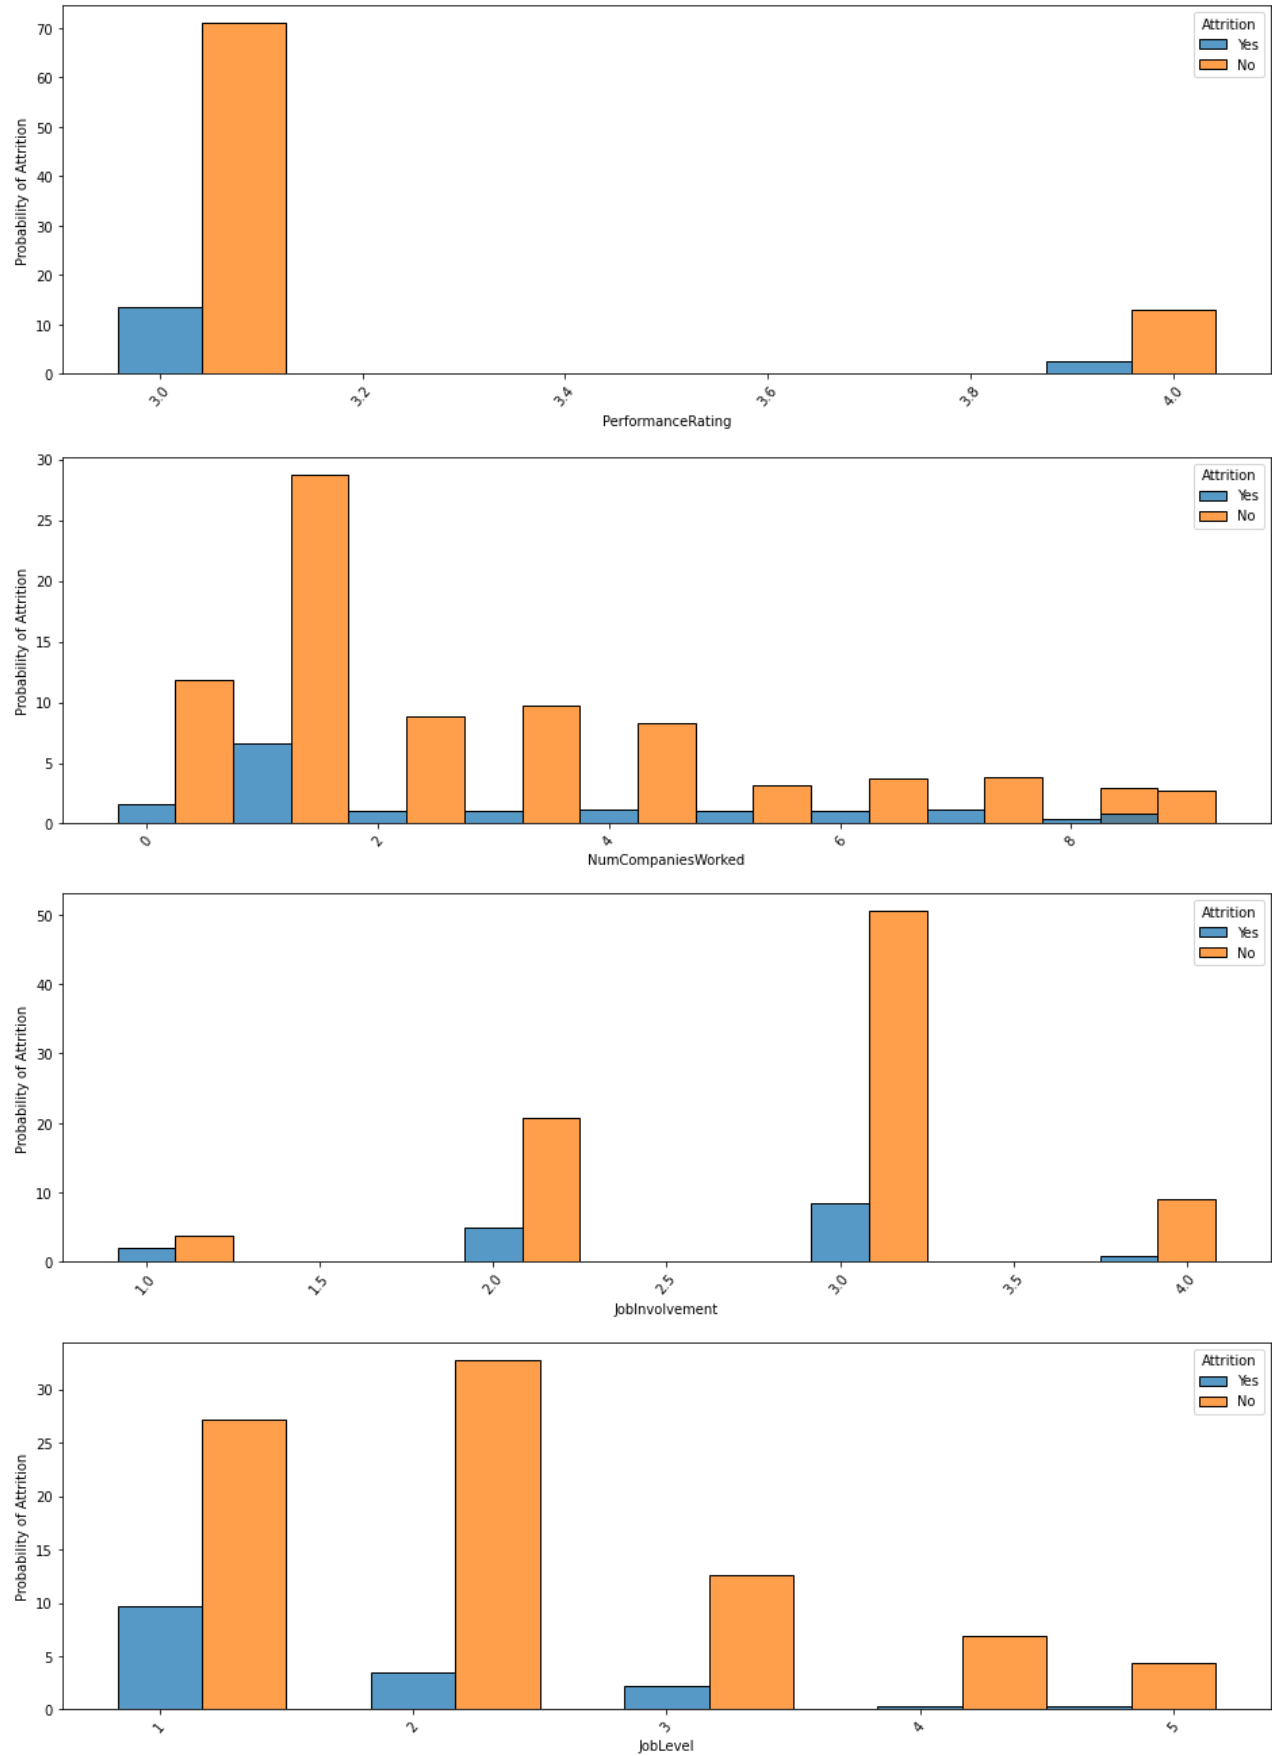
In [11]:

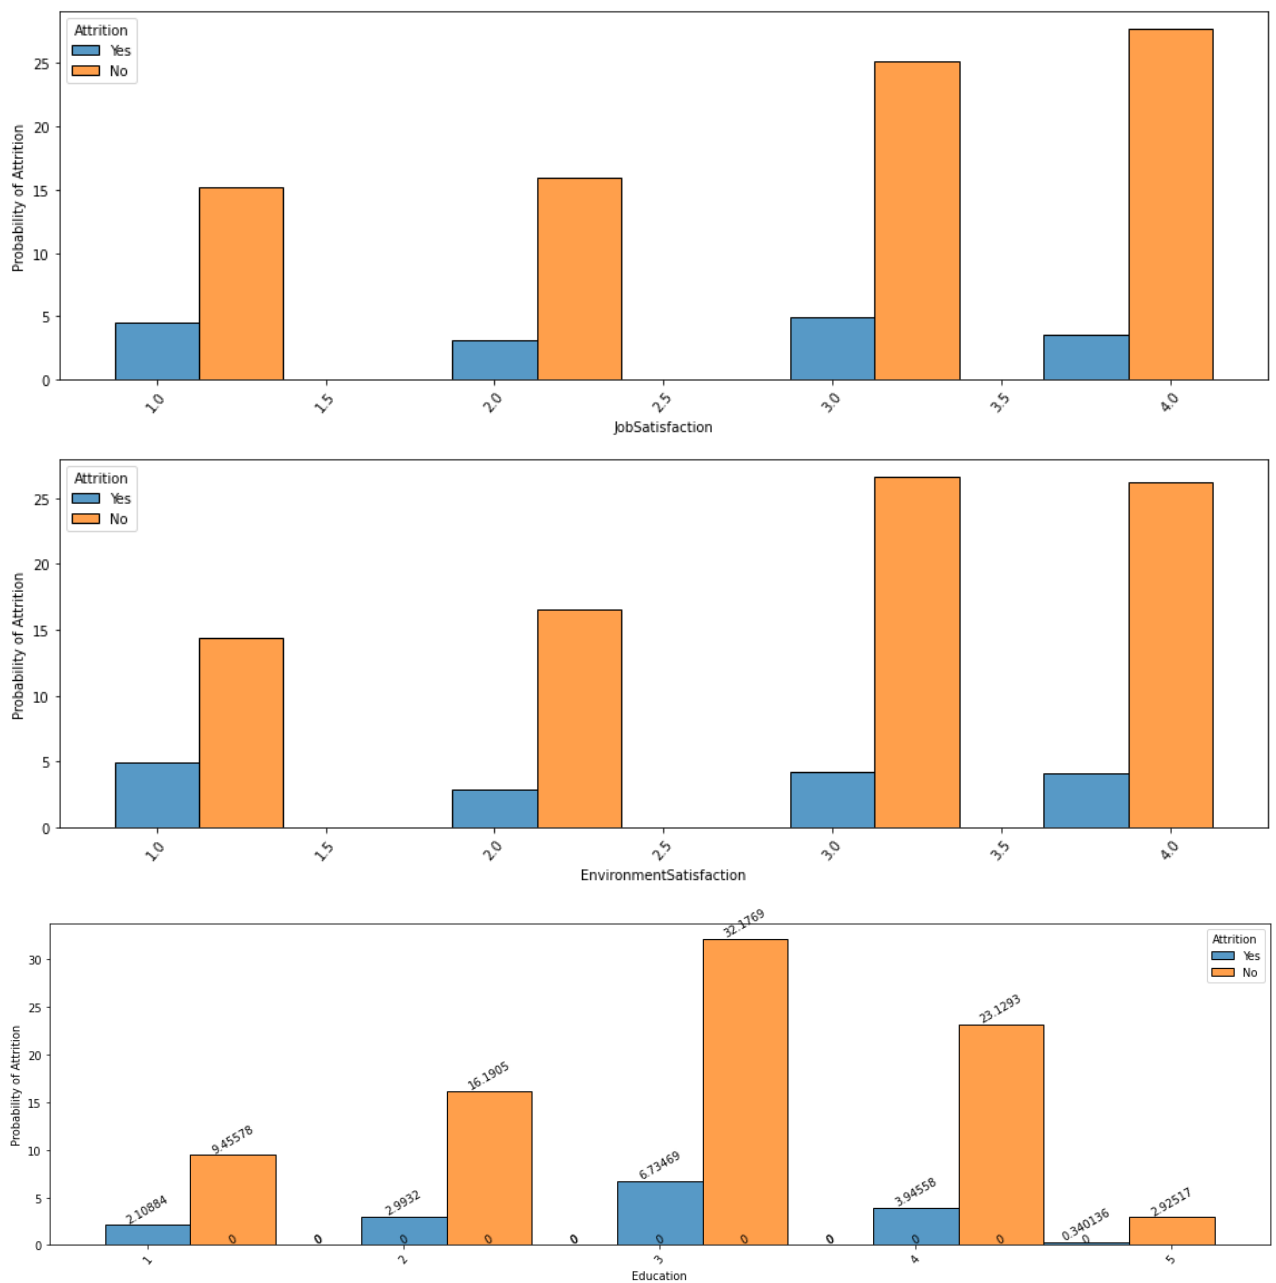
```
cols = ['WorkLifeBalance', 'TrainingTimesLastYear', 'StockOptionLevel',
        'RelationshipSatisfaction', 'PerformanceRating', 'NumCompaniesWorked',
        'JobInvolvement', 'JobLevel', 'JobSatisfaction',
        'EnvironmentSatisfaction', 'Education']

for col in cols:
    fig, axes = plt.subplots(figsize=(16,5))
    sns.histplot(dataset, x=dataset[col], hue="Attrition", stat="percent", multi
    axes.set(ylabel="Probability of Attrition")
    plt.xticks(rotation=50, fontsize=10)
    for cont in axes.containers:
```

```
axes.bar_label(cont,rotation=30,fontsize=10)  
plt.tight_layout()  
plt.show()
```







Step 4: Performance Prediction

```
In [12]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
```

```
In [13]: x.head()
x.columns
```

```
Out[13]: Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education',
              'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
              'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
              'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
              'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
              'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
```

```
'YearsSinceLastPromotion', 'YearsWithCurrManager', 'Attrition_No',
'Attrition_Yes', 'BusinessTravel_Non-Travel',
'BusinessTravel_Travel_Frequently', 'BusinessTravel_Travel_Rarely',
'Department_Human Resources', 'Department_Research & Development',
'Department_Sales', 'EducationField_Human Resources',
'EducationField_Life Sciences', 'EducationField_Marketing',
'EducationField_Medical', 'EducationField_Other',
'EducationField_Technical Degree', 'Gender_Female', 'Gender_Male',
'JobRole_Healthcare Representative', 'JobRole_Human Resources',
'JobRole_Laboratory Technician', 'JobRole_Manager',
'JobRole_Manufacturing Director', 'JobRole_Research Director',
'JobRole_Research Scientist', 'JobRole_Sales Executive',
'JobRole_Sales Representative', 'MaritalStatus_Divorced',
'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
'OverTime_Yes'],
dtype='object')
```

I realize in the workplace, companies may not have easy access to this quantity of features, but to improve accuracy I will use them all.

In [25]:

```
# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(x[['Age', 'DailyRate', 'DistanceFromHome',
'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
'PercentSalaryHike', 'RelationshipSatisfaction',
'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
'YearsSinceLastPromotion', 'YearsWithCurrManager', 'Attrition_No',
'Attrition_Yes', 'BusinessTravel_Non-Travel',
'BusinessTravel_Travel_Frequently', 'BusinessTravel_Travel_Rarely',
'Department_Human Resources', 'Department_Research & Development',
'Department_Sales', 'EducationField_Human Resources',
'EducationField_Life Sciences', 'EducationField_Marketing',
'EducationField_Medical', 'EducationField_Other',
'EducationField_Technical Degree', 'Gender_Female', 'Gender_Male',
'JobRole_Healthcare Representative', 'JobRole_Human Resources',
'JobRole_Laboratory Technician', 'JobRole_Manager',
'JobRole_Manufacturing Director', 'JobRole_Research Director',
'JobRole_Research Scientist', 'JobRole_Sales Executive',
'JobRole_Sales Representative', 'MaritalStatus_Divorced',
'MaritalStatus_Married', 'MaritalStatus_Single', 'OverTime_No',
'OverTime_Yes']], x['PerformanceRating'], test_size=0.25, random_state=42)

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)

print('MAE:', mae)
X_test.head()
```

MAE: 0.17656749614659223

Out[25]:

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInv
1041	28	866	5	3	4	84	

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInv
184	53	1084		13	2	4	57
1222	24	240		22	1	4	58
67	45	1339		7	3	2	59
220	36	1396		5	2	4	62

5 rows × 52 columns

```
In [15]: params = x.drop('PerformanceRating', axis=1)
```

Test first line in dataset, employee 0, and see how our models predicts their performance rating.

```
In [16]: x.head()
```

```
Out[16]:
```

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolve
0	41	1102		1	2	2	94
1	49	279		8	1	3	61
2	37	1373		2	2	4	92
3	33	1392		3	4	4	56
4	27	591		2	1	1	40

5 rows × 53 columns

```
In [17]: employee_data = params.loc[[0]]
prediction = model.predict(employee_data)
print(f'This employee is predicted to receive a performance rating of {prediction[0]}')
```

This employee is predicted to receive a performance rating of 2.8223723789723687

```
In [18]: actual = x.loc[0]['PerformanceRating']
```

The model predicted that this employee would receive a performance score of 2.82, and the employee actually received a score of 3. Lets see if other models could be more accurate.

```
In [19]: from sklearn.tree import DecisionTreeClassifier
```

```
In [20]: # Create a decision tree classifier
dtc = DecisionTreeClassifier()

# Fit the classifier to the training data
dtc.fit(X_train, y_train)

y_pred = dtc.predict(X_test)
```

```
prediction = dtc.predict(employee_data)

# Print the prediction
print(f'This employee is predicted to receive a performance rating of {prediction}')

This employee is predicted to receive a performance rating of [3]
```

```
In [21]: from sklearn.metrics import accuracy_score
```

```
In [22]: accuracy = accuracy_score(y_test, y_pred)
accuracy
```

```
Out[22]: 1.0
```

Decision Tree Classifier has a perfect score. In reality, I would train this model on a smaller, more feasible list of features, so other HR members could easily enter data. See below for an example.

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(x[['Age', 'JobLevel',
'MonthlyIncome', 'TotalWorkingYears', 'TrainingTimesLastYear',
'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
'YearsSinceLastPromotion', 'Gender_Female', 'Gender_Male',
]], x['PerformanceRating'], test_size=0.25, random_state=42)
```

```
In [24]: # Create a decision tree classifier
dtc = DecisionTreeClassifier()

# Fit the classifier to the training data
dtc.fit(X_train, y_train)

y_pred = dtc.predict(X_test)

employee_data = pd.DataFrame({'Age': 40, 'JobLevel':3,
'MonthlyIncome':5000, 'TotalWorkingYears':5, 'TrainingTimesLastYear':0,
'WorkLifeBalance':1, 'YearsAtCompany':5, 'YearsInCurrentRole':2,
'YearsSinceLastPromotion':2, 'Gender_Female':1, 'Gender_Male':0}, index=[0])

prediction = dtc.predict(employee_data.loc[[0]])

# Print the prediction
print(f'This employee is predicted to receive a performance rating of {prediction}')

This employee is predicted to receive a performance rating of [4]
```

```
In [ ]:
```